



 FACULTEIT
INGENIEURSWETENSCHAPPEN



Computer Algebra for Cryptography

Project I

FINAL REPORT

Robin Martens

ACADEMIC YEAR 2024-2025

Task 1

Task 1.e

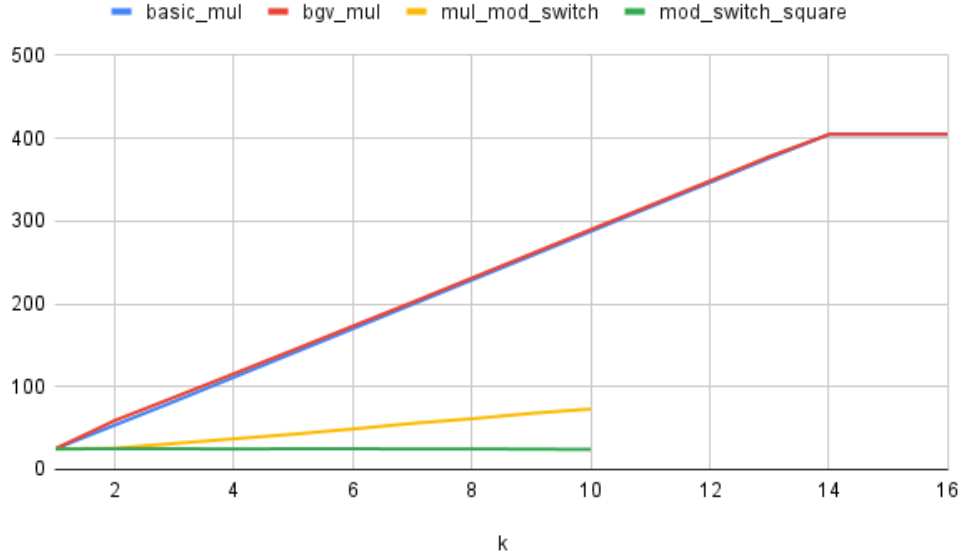


Figure 1: Plot of the different noise bounds for the cases given in the assignment.

Figure 1 shows the noise bound for using just BGVBASICMUL (in blue), for BGVMUL (in red), for BGVMUL, followed by a BGVMODSWITCH (in yellow) and for the repeated squaring (in green). We here clearly see why the modulus switching is necessary to prevent a build up of noise in the ciphertext. Note that the y -axis displays the logarithm of the actual noise bound as defined in the assignment.

We see a similar pattern for both BGVBASICMUL and BGVMUL, which can be explained by the fact that the length of the ciphertext does not influence the noise bound (as we take the infinity norm of the partial decrypt, which is already central reduced). The quadratic noise growth is explained by working out the decryption of the product of two ciphertexts:

$$\begin{aligned}
 \mathcal{D}(\text{ct} \cdot \text{ct}') &= (\mathbf{c}_0 + \mathbf{c}_1 \cdot \mathbf{s}) \cdot (\mathbf{c}'_0 + \mathbf{c}'_1 \cdot \mathbf{s}) \\
 &= (\mathbf{m} + \mathbf{p}\mathbf{e}) \cdot (\mathbf{m}' + \mathbf{p}\mathbf{e}') \\
 &= \mathbf{m}\mathbf{m}' + \mathbf{m}\mathbf{e}' + \mathbf{m}'\mathbf{e} + \mathbf{e}\mathbf{e}'.
 \end{aligned} \tag{1}$$

For BGVMUL, followed by a BGVMODSWITCH, we can clearly see that the noise is reduced to approximately the original level, by using modulus switching.

We get a faulty decryption from the moment that the error term of the decryption $\mathbf{p}\mathbf{e}'$ becomes larger than $\frac{q}{2}$. For the standard parameter set, we get:

$$\log_2 \left(\frac{q}{2} \right) \approx 404.69. \tag{2}$$

When the noise bound comes close to this bound, we get errors in the decryption. From the figure we can derive that we get errors from $k = 14$ and above. A very rough bound can be derived from the fact that we have $e_0 2^k < \frac{q}{2}$. With $\log_2(e_0) \approx 22$, we find $k = 18$, which is higher than we found in the figure, but is explained by the fact that $\mathbf{m}\mathbf{e}' + \mathbf{m}'\mathbf{e}$ also contribute to the noise term.

For multiplying a bunch of ciphertexts, we have to perform the BGVMUL with modulus switching, as long as it's possible based on the level ℓ . Thereafter, the noise will inevitable grow (approximately) quadratically until we reach the noise bound and have a faulty decryption.

Task 4

Task 4.a

The intuition is that the error on the equations cannot become too large w.r.t. the value of the equation itself. If we take the worst case, the maximum value of a single equation is $(q-1)n$ (we can take a more liberal bound if we take into account that $\frac{1}{3}$ of the secret key is supposed to be 0, or even more liberal if take into account the cancellation due to the 1 and -1 in the secret key). The maximum value of an error coefficient is B , so we get the bound

$$(q-1)n > B. \quad (3)$$

Task 4.b

The idea behind the lattice attack is to create a lattice where the shortest vector of the lattice contains the secret key \mathbf{s} . We thus have to create a target vector \mathbf{t} which lies inside the lattice and is much shorter than the expected shortest vector (using the Gaussian + Stirling heuristic) in the lattice.

If we want to create an $(n+3)$ -dimensional lattice, we consider the following matrix (where the rows of the matrix form a basis for the lattice):

$$M = \begin{bmatrix} 1 & \cdots & 0 & 0 & a_1 & 0 & 0 \\ 0 & 1 & \cdots & 0 & a_2 & 0 & 0 \\ \vdots & & \ddots & 0 & \vdots & 0 & 0 \\ 0 & & & 1 & a_n & 0 & 0 \\ 0 & \cdots & \cdots & 0 & p & 1 & 0 \\ 0 & \cdots & \cdots & 0 & B[1] & 0 & 1 \\ 0 & \cdots & \cdots & 0 & q & 0 & 0 \end{bmatrix}. \quad (4)$$

The first n columns are just an identity matrix, while the $(n+1)$ th column corresponds to equation (2) (which is the first row of the circular matrix for the negacyclic convolution, which corresponds to polynomial multiplication modulo $x^N + 1$), where we have expressed the mod q explicitly as a factor kq in the equation.

This way, if we multiply M with $[\mathbf{s} \ E[1] \ -1 \ k]$, we get:

$$[\mathbf{s} \ E[1] \ -1 \ k] \cdot M = [\mathbf{s} \ 0 \ E[1] \ 0]. \quad (5)$$

For the volume we have, as $d = n$:

$$\det M = q \det(I_{(n+2) \times (n+2)}) = q, \quad (6)$$

from which we can calculate the expected length of the shortest vector $\lambda_1(L)$:

$$\begin{aligned} \lambda_1(L) &\approx \sqrt{\frac{d}{2\pi e}} \text{vol}(L)^{1/d} \\ &\approx \sqrt{\frac{N+3}{2\pi e}} q^{\frac{k}{N+3}}, \end{aligned} \quad (7)$$

if we take q to be of the form q^k . The expected length of \mathbf{t} is (worst case):

$$\|\mathbf{t}\| \approx \sqrt{\frac{2}{3}N + B^2}, \quad (8)$$

as \mathbf{s} is taken uniformly and is ternary, so $\frac{2}{3}$ of the elements are either 1 or -1 . If we combine the above two equations and solve for k , we find:

$$k \geq 4.07. \quad (9)$$

If we wouldn't take out p explicitly from $E[1]$, we could not put p in the matrix as above and we cannot put $pE[1]$ in the target vector, as it would be too large (or we would have to take k very large).

Task 4.d

Experimentally, we find that we find the correct vector for $k \geq 5$, which corresponds to the bound for k found in the previous subtask.

Task 5**Task 5.a**

For $k = 2$, the decryption equation becomes

$$\mathbf{r} = [\mathbf{ct}[0] + \mathbf{ct}[1]\mathbf{s}]_q \mod p. \quad (10)$$

As \mathbf{s} is small, it will not be central reduced, nor be reduced mod p (if $p > 2$), such that if we plug in $\mathbf{ct} = [0, 1]$, we get $\mathbf{r} = \mathbf{s}$.

If $p = 2$, however, the mod p , *will* reduce \mathbf{s} (in this case, as \mathbf{s} is ternary, all the -1 will be reduced to 1).

Task 5.c

As explained above, there will be a decryption error when the noise term of the partial decryption is greater than $\frac{q}{2}$, as then the noise wraps around due to the centered reduction. Now, by construction we have $q = 1 \mod p$ and thus $(q - 1) = 0 \mod p$ and also $\frac{q-1}{2} = 0 \mod p$. Secondly, we note that

$$\begin{cases} \frac{q-1}{2} < \frac{q}{2} \\ \frac{q-1}{2} + 1 > \frac{q}{2} \end{cases} \quad (11)$$

This means that if we create a ciphertext with noise bound $\frac{q-1}{2}$, it will be a borderline case. Thirdly, we notice that the noise term is always a multiple of p and that by adding multiples of p to the any part of the ciphertext will not change the decryption as long as we don't surpass the noise bound.

We start with the public key as a valid ciphertext, which is defined as

$$\mathbf{b} = \mathbf{a} \cdot \mathbf{s} + p\mathbf{e}, \quad (12)$$

where the coefficients of \mathbf{e} are in the interval $[-B, B]$, with B small. The fact that B is small can be exploited as follows. Suppose that the i th coefficient of \mathbf{e} is $e_i \in [-B, B]$. Now we add $\frac{q-1}{2}$ to e_i , such that we either have an overshoot or undershoot of the noise bound (depending on the fact of the original e_i was positive or negative). Note that in the special case that $e_i = 0$, we are already at the noise bound. In general, we then loop twice over the interval $[1, B]$. In the first round, we subtract p from e_i (to test for overshoot) and in the second round, we add p to e_i (to test for undershoot). Somewhere in this looping, we reach the bound $\frac{q-1}{2}$, which indicates that we found e_i .

The final things which we have to do is finding a way to check if we reached the bound $\frac{q-1}{2}$, as we do not have access to the `BGVPartialDecrypt` method. If some $e_i = \frac{q-1}{2}$, the decryption will still be correct, but if we add 1 to e_i , we get (taking into account the centered reduction):

$$\frac{q-1}{2} + 1 - q = \frac{1-q}{2} = -\frac{q-1}{2}, \quad (13)$$

which is *still* a multiple of p , such that it disappears modulo p . This indicates that we still get the original message that we encrypted, which is a decryption failure, as we should get $\mathbf{m} + 1$, following the decryption equation. It suffices thus to check that `BGVDecrypt(ct, sk) eq BGVDecrypt(<[ct[1][1]+1, ct[1][2]], ct[2]>, sk)`.

This can also be generalized to work for all coefficients at the same time. Then, the followings steps should be adapted:

- Instead of adding $\frac{q-1}{2}$ in the first step, we add the polynomial of degree $N - 1$ with all coefficients equal to $\frac{q-1}{2}$ (i.e. $\sum_{n=0}^{N-1} \left(\frac{q-1}{2}\right) x^i$).

- In the iterations, we do a similar change: instead of adding p , we add the polynomial of degree $N - 1$ with all coefficients equal to p .
- To check which coefficients already reached the threshold, we add a polynomial of degree $N - 1$ with all coefficients 1 and decrypt. We then check which coefficients are equal (using the principle described above) to the decryption of the ciphertext before we added the polynomial.

The methods where we find all coefficients at once is *much* more efficient then the algorithm where all coefficients are found one by one (experimentally up to a factor 40 less).

Task 6

Task 6.c

Yes, if we want to multiply polynomials over the ring $\mathbb{Z}_q[x] \setminus (x^N - 1)$, we have to use the so-called positive wrapped convolution or cyclic convolution and for working over the ring $\mathbb{Z}_q[x] \setminus (x^N + 1)$, we have to use the so-called negative wrapped convolution (or negacyclic convolution). The main point is that by evaluating the NTT in the roots of $(x^N - 1)$, we're working in the above mentioned ring. We can use the same principle to work over the ring module $(x^N + 1)$, namely by evaluating the NTT in the roots $(x^N + 1)$ instead of the roots of $(x^N - 1)$. We call these roots ψ , the $2n$ -th root of unity such that $\psi^2 = \omega$ in \mathbb{Z}_q , as $x^{2n} - 1 = (x^n - 1)(x^n + 1)$. We can then evaluate the NTT in the odd powers of ψ , which gives $[1, 2]$:

$$\begin{aligned} \text{NTT}_{\psi}(a)_j &= \sum_{i=0}^{n-1} a_i \psi_{2n}^{i(2j+1)} \\ &= \sum_{i=0}^{n-1} a_i \omega_n^{ij} \psi_{2n}^i \\ &= \text{NTT}_{\omega}([\psi_{2n} a_0, \dots, \psi_{2n}^{n-1} a_{n-1}]) \end{aligned} \tag{14}$$

Task 6.d

Yes, we could use the NTT and the INTT to do the encoding. The main idea of encoding is to perform multiplication and addition in R_p , which is then a point-wise multiplication in \mathbb{F}_p . As we want to have pointwise multiplication over a sequence in \mathbb{F}_p , which corresponds to multiplication in R_p , we take the inverse NTT for encoding and the NTT for decoding.

References

- [1] *Number Theoretic Transform - A Gentle Introduction: Part II*. June 2023. URL: <https://cryptographycaffe.sandboxaq.com/posts/ntt-02/>.
- [2] Ardianto Satriawan et al. “Conceptual review on number theoretic transform and comprehensive review on its implementations”. In: *IEEE Access* 11 (Jan. 2023), pp. 70288–70316. DOI: 10.1109/access.2023.3294446. URL: <https://doi.org/10.1109/access.2023.3294446>.